In [1]:
```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing,svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [2]:
```
df=pd.read_csv(r"C:\Users\pappu\Downloads\fiat500_VehicleSelection_Dataset.csv
df
```

Out[2]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611560
1	2	pop	51	1186	32500	1	45.666359	12.241890
2	3	sport	74	4658	142228	1	45.503300	11.417840
3	4	lounge	51	2739	160000	1	40.633171	17.634609
4	5	pop	73	3074	106880	1	41.903221	12.495650
...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870
1535	1536	pop	51	2223	60457	1	45.481541	9.413480
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270
1537	1538	pop	51	1766	54276	1	40.323410	17.568270

1538 rows × 9 columns

In [3]:
```
df=df[['price','age_in_days']]
df.columns=['price','age']
```

In [4]:
```
df.head(5)
```

Out[4]:

	price	age
0	8900	882
1	8800	1186
2	4200	4658
3	6000	2739
4	5700	3074

In [5]:
```
df.describe()
```

Out[5]:

	price	age
count	1538.000000	1538.000000
mean	8576.003901	1650.980494
std	1939.958641	1289.522278
min	2500.000000	366.000000
25%	7122.500000	670.000000
50%	9000.000000	1035.000000
75%	10000.000000	2616.000000
max	11100.000000	4658.000000

In [6]:
```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  ------  --------------  -----
 0   price   1538 non-null   int64
 1   age     1538 non-null   int64
dtypes: int64(2)
memory usage: 24.2 KB
```

In [7]:
```
df.isna().any()
```

Out[7]:
```
price    False
age      False
dtype: bool
```

In [8]:
```
sns.lmplot(x='price',y='age',data=df,order=2,ci=None)
```

Out[8]: <seaborn.axisgrid.FacetGrid at 0x2ad11dc46a0>

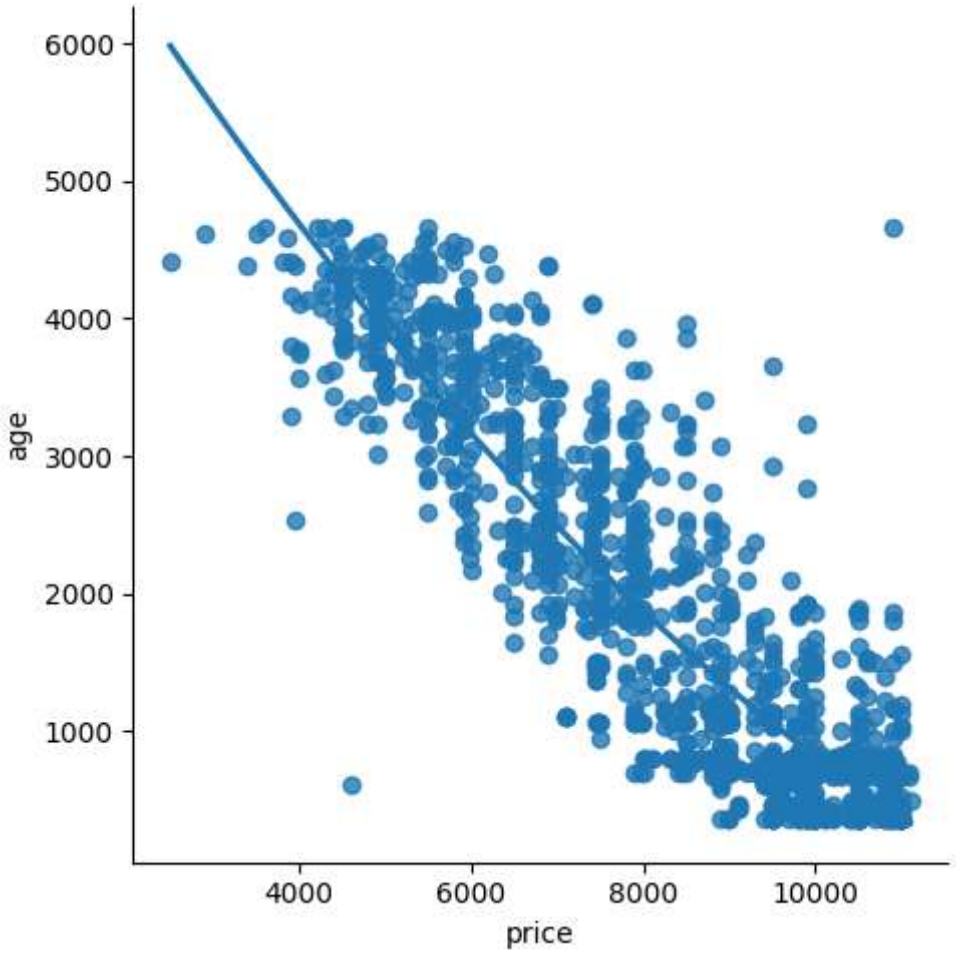

In [9]:
```
x=np.array(df['price']).reshape(-1,1)
y=np.array(df['age']).reshape(-1,1)
```

In [11]:
```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```

```
0.7766219113040318
```

In [12]:
```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

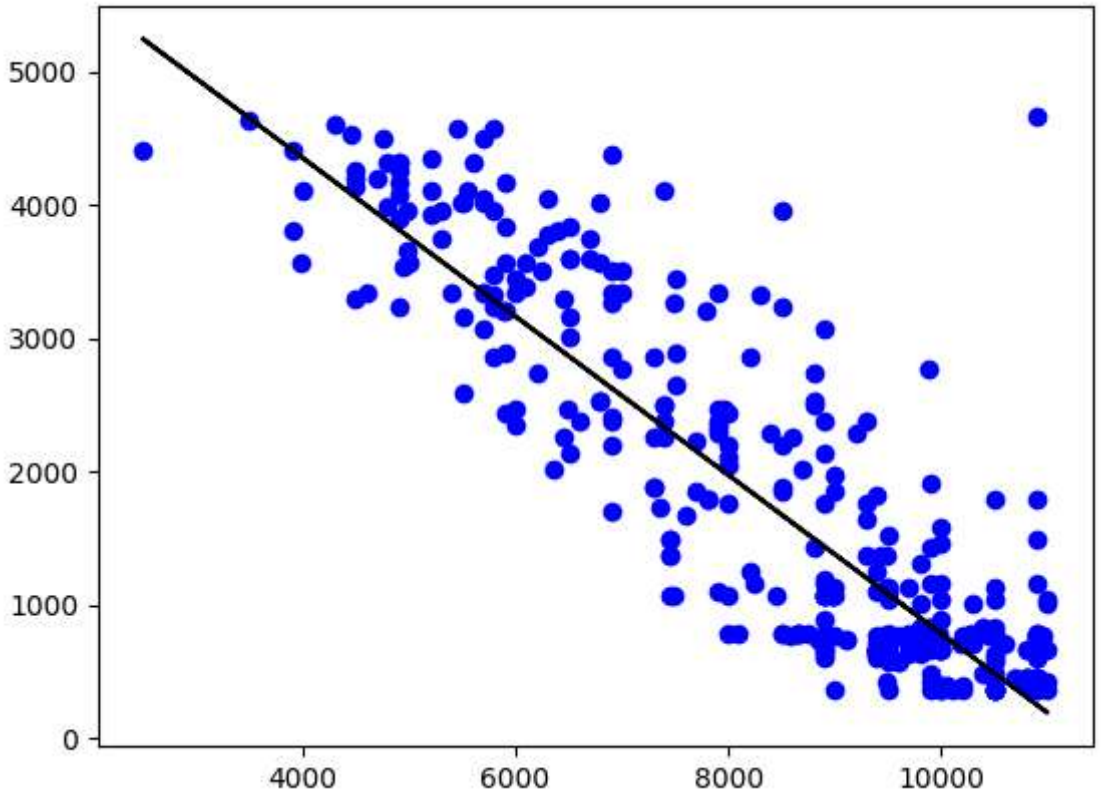

In [16]:
```
df500=df[:][:500]
sns.lmplot(x="price",y="age",data=df500,order=1,ci=None)
```

Out[16]: <seaborn.axisgrid.FacetGrid at 0x2ad1767f100>

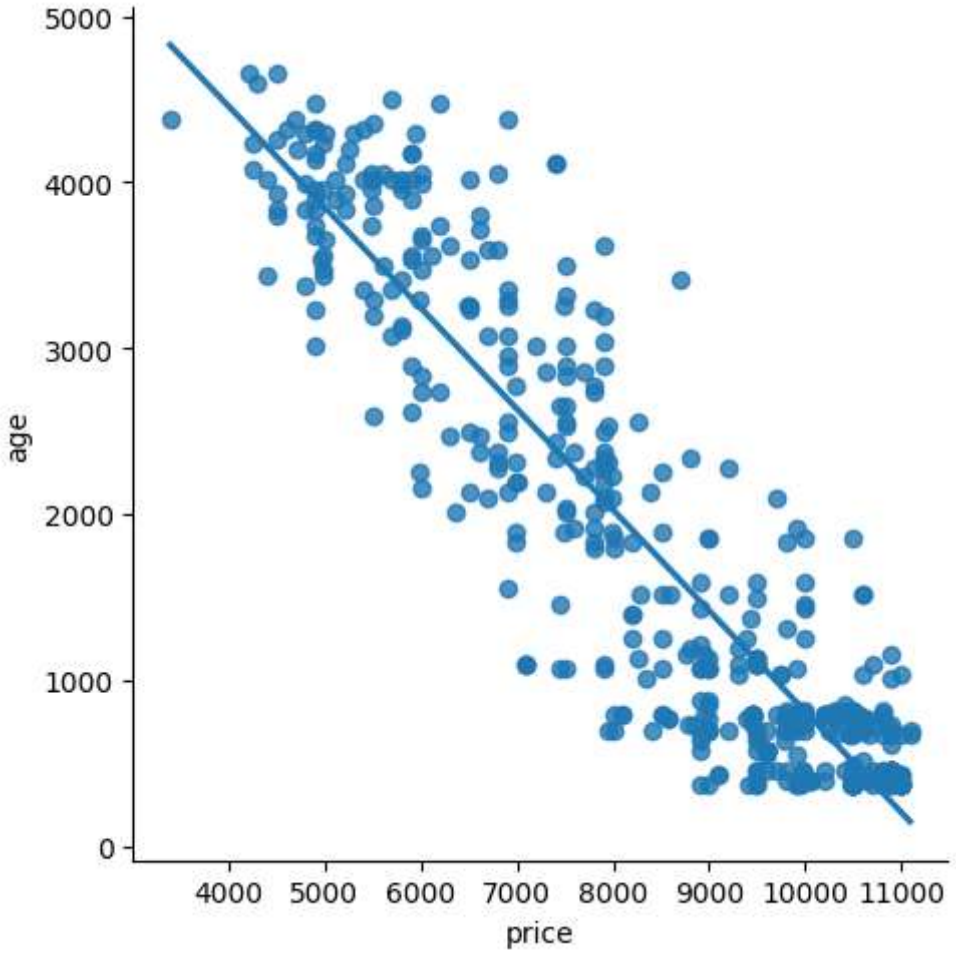

In [17]:
```
x=np.array(df500['price']).reshape(-1,1)
y=np.array(df500['age']).reshape(-1,1)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

```
0.823729400746001
```

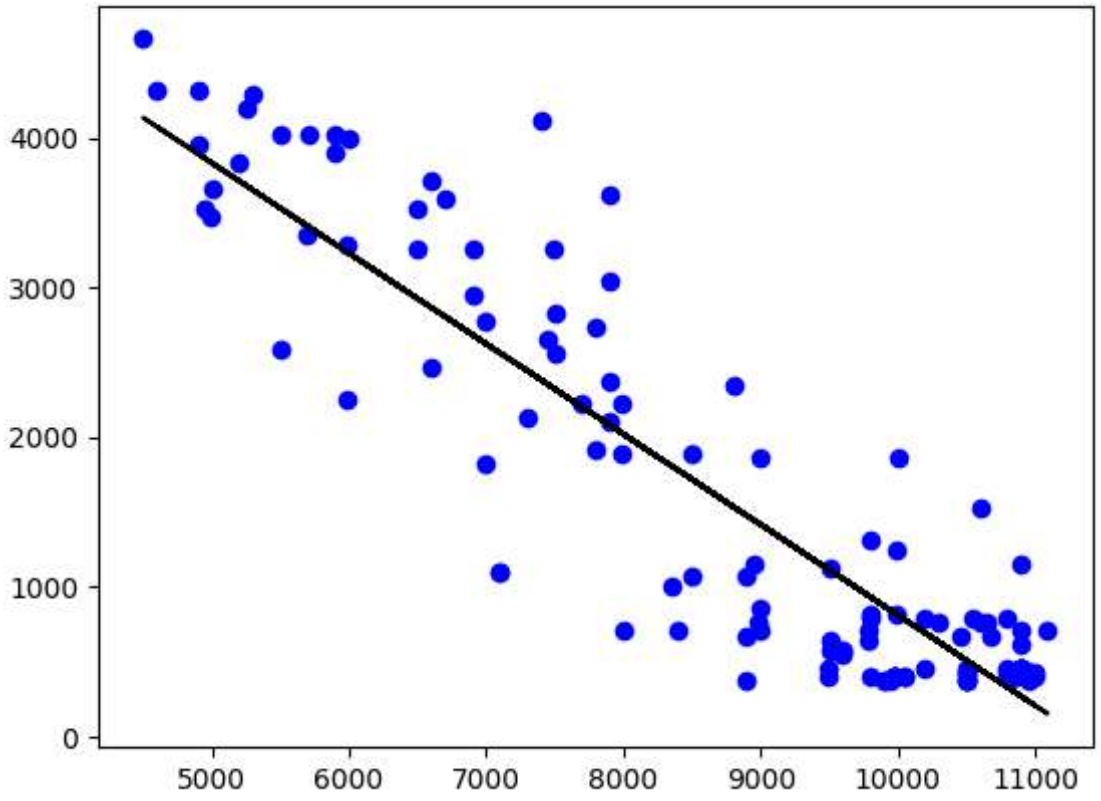

In [20]:
```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

```
R2 score: 0.823729400746001
```

In [ ]: